

## Submission Assignment #2 (word2vec (43 Points))

Instructor: Christopher Manning

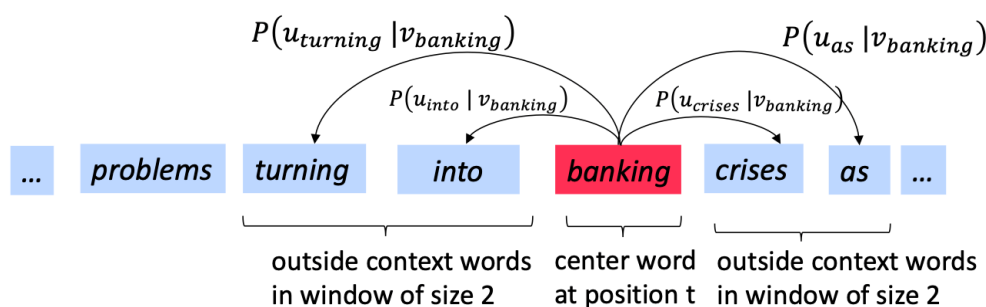
Name: Jianpan Gun

## Problem 1: Understanding word2vec

(3+5+5+3+4+3=23 points)

Let's have a quick refresher on the **word2vec** algorithm. The key insight behind **word2vec** is that 'a word is known by the company it keeps'. Concretely, suppose we have a 'center' word  $c$  and a contextual window surrounding  $c$ . We shall refer to words that lie in this contextual window as 'outside words'. For example, in Figure 1 we see that the center word  $c$  is 'banking'. Since the context window size is 2, the outside words are 'turning', 'into', 'crises', and 'as'.

The goal of the skip-gram **word2vec** algorithm is to accurately learn the probability distribution  $P(O|C)$ . Given a specific word  $o$  and a specific word  $c$ , we want to calculate  $P(O = o|C = c)$ , which is the probability that word  $o$  is an 'outside' word for  $c$ , i.e., the probability that  $o$  falls within the contextual window of  $c$ .

Figure 1: The **word2vec** skip-gram prediction model with window size 2.

In **word2vec**, the conditional probability distribution is given by taking vector dot-products and applying the softmax function:

$$P(O = o|C = c) = \frac{\exp(\mathbf{u}_o^\top \mathbf{v}_c)}{\sum_{w \in V_{\text{ocab}}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} \quad (1)$$

Here,  $\mathbf{u}_o$  is the 'outside' vector representing outside word  $o$ , and  $\mathbf{v}_c$  is the 'center' vector representing center word  $c$ . To contain these parameters, we have two matrices,  $\mathbf{U}$  and  $\mathbf{V}$ . The columns of  $\mathbf{U}$  are all the 'outside' vectors  $\mathbf{u}_w$ . The columns of  $\mathbf{V}$  are all of the 'center' vectors  $\mathbf{v}_w$ . Both  $\mathbf{U}$  and  $\mathbf{V}$  contain a vector for every  $w \in \text{Vocabulary}$ .

Recall from lectures that, for a single pair of words  $c$  and  $o$ , the loss is given by:

$$J_{\text{naive-softmax}}(\mathbf{v}_c, o, \mathbf{U}) = -\log P(O = o|C = c) \quad (2)$$

Another way to view this loss is as the cross-entropy between the true distribution  $\mathbf{y}$  and the predicted distribution  $\hat{\mathbf{y}}$ . Here, both  $\mathbf{y}$  and  $\hat{\mathbf{y}}$  are vectors with length equal to the number of words in the vocabulary. Furthermore, the  $k^{\text{th}}$  entry in these vectors indicates the conditional probability of the  $k^{\text{th}}$  word being an 'outside word' for the given  $c$ . The true empirical distribution  $\mathbf{y}$  is a one-hot vector with a 1 for the true outside word  $o$ , and 0 everywhere else. The predicted distribution  $\hat{\mathbf{y}}$  is the probability distribution  $P(O|C = c)$  given by our model in equation (1).

(a) (3 points) Show that the naive-softmax loss given in Equation (2) is the same as the cross-entropy loss between  $\mathbf{y}$  and  $\hat{\mathbf{y}}$ ; i.e., show that  $-\sum_{w \in V_{\text{ocab}}} y_w \log(\hat{y}_w) = -\log(\hat{y}_o)$

A: Due to the  $\mathbf{y}_i$  is a one-hot vector, that is,  $\mathbb{I}_i$ , which means that only have  $i$ -th dim is 1. So,

$$\mathcal{L}_{ce} = - \sum_{w \in V_{ocab}} y_w \log(\hat{y}_w) = -\mathbf{y}_o \log(\hat{\mathbf{y}}_o) = -\log(\hat{y}_o) \quad (3)$$

(b) (5 points) Compute the partial derivative of  $J_{\text{naive-softmax}}(\mathbf{v}_c, o, U)$  with respect to  $\mathbf{v}_c$ . Please write your answer in terms of  $\mathbf{y}$ ,  $\hat{\mathbf{y}}$ , and  $U$

$$\begin{aligned} \frac{\partial J_{\text{naive-softmax}}}{\partial \mathbf{v}_c}(\mathbf{v}_c, o, U) &= \frac{\partial}{\partial \mathbf{v}_c} -\log \frac{\exp(\mathbf{u}_o^\top \mathbf{v}_c)}{\sum_{w \in V_{ocab}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} = -\mathbf{u}_o + \frac{\partial}{\partial \mathbf{v}_c} \log \sum_{w \in V_{ocab}} \exp(\mathbf{u}_w^\top \mathbf{v}_c) \\ &= -\mathbf{u}_o + \frac{1}{\sum_{w \in V_{ocab}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} \sum_{w \in V_{ocab}} \exp(\mathbf{u}_w^\top \mathbf{v}_c) \mathbf{u}_w \\ &= -\mathbf{y}^\top U + \sum_{w \in V_{ocab}} P(O = w | C = c) \mathbf{u}_w \\ &= -\mathbf{y}^\top U + \hat{\mathbf{y}}^\top U = (\hat{\mathbf{y}} - \mathbf{y})^\top U \end{aligned} \quad (4)$$

(c) (5 points) Compute the partial derivatives of  $J_{\text{naive-softmax}}(\mathbf{v}_c, o, U)$  with respect to each of the ‘outside’ word vectors,  $\mathbf{u}_w$ ’s. There will be two cases: when  $w = o$ , the true ‘outside’ word vector; and  $w \neq o$ , for all other words. Please write your answer in terms of  $\mathbf{y}$ ,  $\hat{\mathbf{y}}$ , and  $\mathbf{v}_c$ .

$$\frac{\partial J_{\text{naive-softmax}}}{\partial \mathbf{u}_w}(\mathbf{v}_c, o, U) = \frac{\partial}{\partial \mathbf{u}_w} -\log \frac{\exp(\mathbf{u}_o^\top \mathbf{v}_c)}{\sum_{w \in V_{ocab}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} = -\frac{\partial \mathbf{u}_o^\top \mathbf{v}_c}{\partial \mathbf{u}_w} + \frac{\partial}{\partial \mathbf{u}_w} \log \sum_{w \in V_{ocab}} \exp(\mathbf{u}_w^\top \mathbf{v}_c) \quad (5)$$

If  $w = o$ :

$$= -\mathbf{v}_c + \frac{\exp(\mathbf{u}_o^\top \mathbf{v}_c) \mathbf{v}_c}{\sum_{w \in V_{ocab}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} = (P(O = o | C = c) - 1) \mathbf{v}_c \quad (6)$$

If  $w \neq o$ :

$$= \frac{\exp(\mathbf{u}_w^\top \mathbf{v}_c) \mathbf{v}_c}{\sum_{w \in V_{ocab}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} = P(O = w | C = c) \mathbf{v}_c \quad (7)$$

So,

$$= (\hat{\mathbf{y}} - \mathbf{y})^\top \mathbf{v}_c \quad (8)$$

(d) (3 Points) The sigmoid function is given by Equation 9:

$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1} \quad (9)$$

Please compute the derivative of  $\sigma(x)$  with respect to  $x$ , where  $x$  is a scalar. Hint: you may want to write your answer in terms of  $\sigma(x)$ .

$$\frac{\partial \sigma(x)}{\partial x} = \frac{\partial}{\partial x} \frac{e^x}{e^x + 1} = \frac{e^x(e^x + 1) - e^x(e^x + 1)'}{(e^x + 1)^2} = \frac{e^{2x}}{(e^x + 1)^2} = \sigma(x)(1 - \sigma(x)) \quad (10)$$

(e) (4 points) Now we shall consider the Negative Sampling loss, which is an alternative to the Naive Softmax loss. Assume that  $K$  negative samples (words) are drawn from the vocabulary. For simplicity of notation we shall refer to them as  $w_1, w_2, \dots, w_K$  and their outside vectors as  $\mathbf{u}_1, \dots, \mathbf{u}_K$ . Note that  $o \notin w_1, \dots, w_K$ . For a center word  $c$  and an outside word  $o$ , the negative sampling loss function is given by:

$$J_{\text{neg-sample}}(\mathbf{v}_c, o, U) = -\log(\sigma(\mathbf{u}_o^\top \mathbf{v}_c)) - \sum_{k=1}^K \log(\sigma(-\mathbf{u}_k^\top \mathbf{v}_c)) \quad (11)$$

for a sample  $w_1, w_2, \dots, w_K$ , where  $\sigma(\cdot)$  is the sigmoid function.

Please repeat parts (b) and (c), computing the partial derivatives of  $J_{\text{neg-sample}}$  with respect to  $v_c$ , with respect to  $u_o$ , and with respect to a negative sample  $u_k$ . Please write your answers in terms of the vectors  $u_o, v_c$ , and  $u_k$ , where  $k \in [1, K]$ . After you've done this, describe with one sentence why this loss function is much more efficient to compute than the naive-softmax loss. Note, you should be able to use your solution to part (d) to help compute the necessary gradients here.

$$\begin{aligned} \frac{\partial J_{\text{neg-sample}}}{\partial v_c}(v_c, o, U) &= \frac{\partial}{\partial v_c} [-\log(\sigma(u_o^\top v_c)) - \sum_{k=1}^K \log(\sigma(-u_k^\top v_c))] = -\frac{\sigma(u_o^\top v_c)(1 - \sigma(u_o^\top v_c))u_o}{\sigma(u_o^\top v_c)} \\ &\quad + \sum_{k=1}^K (1 - \sigma(-u_k^\top v_c))u_k = -(1 - \sigma(u_o^\top v_c))u_o + \sum_{k=1}^K (1 - \sigma(-u_k^\top v_c))u_k \end{aligned} \quad (12)$$

$$\begin{aligned} \frac{\partial J_{\text{neg-sample}}}{\partial u_o}(v_c, o, U) &= \frac{\partial}{\partial u_o} [-\log(\sigma(u_o^\top v_c)) - \sum_{k=1}^K \log(\sigma(-u_k^\top v_c))] = -\frac{\sigma(u_o^\top v_c)(1 - \sigma(u_o^\top v_c))v_c}{\sigma(u_o^\top v_c)} \\ &= -(1 - \sigma(u_o^\top v_c))v_c \end{aligned} \quad (13)$$

$$\frac{\partial J_{\text{neg-sample}}}{\partial u_k}(v_c, o, U) = \frac{\partial}{\partial u_k} [-\log(\sigma(u_o^\top v_c)) - \sum_{k=1}^K \log(\sigma(-u_k^\top v_c))] = (1 - \sigma(-u_k^\top v_c))v_c \quad (14)$$

(f) (3 points) Suppose the center word is  $c = wt$  and the context window is  $w_{t-m}, \dots, w_{t-1}, w_t, w_{t+1}, \dots, w_{t+m}$ , where  $m$  is the context window size. Recall that for the skip-gram version of word2vec, the total loss for the context window is:

$$J_{\text{skip-gram}}(v_c, w_{t-m}, \dots, w_{t+m}, U) = \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} J(v_c, w_{t+j}, U) \quad (15)$$

Here,  $J(v_c, w_{t+j}, U)$  represents an arbitrary loss term for the center word  $c = wt$  and outside word  $w_{t+j}$ .  $J(v_c, w_{t+j}, U)$  could be  $J_{\text{neg-sample}}(v_c, w_{t+j}, U)$  or  $J_{\text{naive-softmax}}(v_c, w_{t+j}, U)$ , depending on your implementation.

Write down three partial derivatives:

- (I)  $\partial J_{\text{skip-gram}}(v_c, w_{t-m}, \dots, w_{t+m}, U) / \partial U$
- (II)  $\partial J_{\text{skip-gram}}(v_c, w_{t-m}, \dots, w_{t+m}, U) / \partial v_c$
- (III)  $\partial J_{\text{skip-gram}}(v_c, w_{t-m}, \dots, w_{t+m}, U) / \partial v_w$ , when  $w \neq c$

Write your answers in terms of  $\partial J(v_c, w_{t+j}, U) / \partial U$  and  $\partial J(v_c, w_{t+j}, U) / \partial v_c$ . This is very simple. Each solution should be one line. Once you're done: Given that you computed the derivatives of  $J(v_c, w_{t+j}, U)$  with respect to all the model parameters  $U$  and  $V$  in parts (a) to (c), you have now computed the derivatives of the full loss function  $J_{\text{skip-gram}}$  with respect to all parameters. You're ready to implement word2vec!

(I)

For native-softmax:

$$\frac{\partial J_{\text{skip-gram}}}{\partial U}(v_c, o, U) = \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} (\hat{y} - y)^\top v_j \quad (16)$$

For neg-sample:

$$\frac{\partial J_{\text{skip-gram}}}{\partial U}(v_c, o, U) = \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} (1 - \sigma(u_o^\top v_j))v_j \quad (17)$$

(II)

For native-softmax:

$$\frac{\partial J_{\text{skip-gram}}}{\partial v_c}(v_c, o, U) = \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} (\hat{y} - y)^\top U \quad (18)$$

For neg-sample:

$$\frac{\partial J_{\text{skip-gram}}}{\partial \mathbf{v}_c}(\mathbf{v}_c, o, U) = \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} (-(1 - \sigma(\mathbf{u}_o^\top \mathbf{v}_j)) \mathbf{u}_o + \sum_{k=1}^K (1 - \sigma(-\mathbf{u}_k^\top \mathbf{v}_j)) \mathbf{u}_k) \quad (19)$$

(III)

$$\frac{\partial J_{\text{skip-gram}}}{\partial \mathbf{v}_w}(\mathbf{v}_c, o, U) = 0 \quad (20)$$

**Problem Implementing word2vec (20 points)**

(12+4+4=20 points)

(a) (12 points) We will start by implementing methods in word2vec.py. First, implement the sigmoid method, which takes in a vector and applies the sigmoid function to it. Then implement the softmax loss and gradient in the naiveSoftmaxLossAndGradient method, and negative sampling loss and gradient in the negSamplingLossAndGradient method. Finally, fill in the implementation for the skip-gram model in the skipgram method. When you are done, test your implementation by running python word2vec.py

Listing 1: word2vec.py function.

```

1 def sigmoid(x):
2     # \sigma(x)=\frac{1}{1+e^{-x}}=\frac{e^x}{e^x+1}
3     s = 1 / (1 + np.exp(-x))
4     return s
5
6 def naiveSoftmaxLossAndGradient(centerWordVec, outsideWordIdx, outsideVectors, ←
dataset):
7     y_hat = softmax(np.dot(centerWordVec, outsideVectors.T))
8
9     loss = -np.log(y_hat[outsideWordIdx]) # ref to equation(2)
10    y = np.zeros_like(y_hat)
11    y[outsideWordIdx] = 1
12    gradCenterVec = np.dot((y_hat - y).T, outsideVectors) # ref to equation(4)
13    gradOutsideVecs = np.dot(
14        (y_hat - y).reshape(-1, 1), centerWordVec.reshape(1, -1)
15    ) # (T, 1) * (1, H) ref to equation(8)
16
17    return loss, gradCenterVec, gradOutsideVecs
18
19 def negSamplingLossAndGradient(
20     centerWordVec, outsideWordIdx, outsideVectors, dataset, K=10
21 ):
22     negSampleWordIndices = getNegativeSamples(outsideWordIdx, dataset, K)
23     indices = [outsideWordIdx] + negSampleWordIndices
24
25     negSampleWordIndices = np.array(negSampleWordIndices)
26
27     pos = sigmoid(np.dot(outsideVectors[outsideWordIdx], centerWordVec))
28     neg = sigmoid(-np.dot(outsideVectors[negSampleWordIndices], centerWordVec))
29     loss = -np.log(pos) - np.log(neg).sum() # ref to equation(11)
30     gradCenterVec = -np.dot((1 - pos), outsideVectors[outsideWordIdx]) + np.dot(
31         1 - neg, outsideVectors[negSampleWordIndices]
32     ) # ref to equation(12)
33     gradOutsideVecs = np.zeros_like(outsideVectors)
34     gradOutsideVecs[outsideWordIdx] = - np.dot(1 - pos, centerWordVec) # ref to ←
equation(13)
35     gradOutside = np.dot((1 - neg).reshape(-1, 1), centerWordVec.reshape(1, -1))
36     for idx, neg_idx in enumerate(negSampleWordIndices):
37         gradOutsideVecs[neg_idx] += gradOutside[idx] # ref to equation(14)

```

```

38
39     return loss, gradCenterVec, gradOutsideVecs
40
41 def skipgram(
42     currentCenterWord,
43     windowSize,
44     outsideWords,
45     word2Ind,
46     centerWordVectors,
47     outsideVectors,
48     dataset,
49     word2vecLossAndGradient=naiveSoftmaxLossAndGradient,
50 ):
51     loss = 0.0
52     gradCenterVecs = np.zeros(centerWordVectors.shape)
53     gradOutsideVectors = np.zeros(outsideVectors.shape)
54
55     centerWordIdx = word2Ind[currentCenterWord]
56     centerWordVector = centerWordVectors[centerWordIdx]
57     for outsideWord in outsideWords:
58         outsideWordIdx = word2Ind[outsideWord]
59         tmp_loss, gradCenterVec, gradOutsideVector = word2vecLossAndGradient(
60             centerWordVector, outsideWordIdx, outsideVectors, dataset
61         )
62         loss += tmp_loss
63         gradCenterVecs[centerWordIdx, :] += gradCenterVec # ref to equation(18,19)
64         gradOutsideVectors += gradOutsideVector
65
66     return loss, gradCenterVecs, gradOutsideVectors

```

---

(b) (4 points) Complete the implementation for your SGD optimizer in the `sgd` method of `sgd.py`. Test your implementation by running `python sgd.py`.

Listing 2: `sgd.py` function.

```

1
2 def sgd(f, x0, step, iterations, postprocessing=None, useSaved=False, PRINT_EVERY↵
   =10):
3     # Anneal learning rate every several iterations
4     ANNEAL_EVERY = 20000
5
6     if useSaved:
7         start_iter, oldx, state = load_saved_params()
8         if start_iter > 0:
9             x0 = oldx
10            step *= 0.5 ** (start_iter / ANNEAL_EVERY)
11
12            if state:
13                random.setstate(state)
14        else:
15            start_iter = 0
16
17        x = x0
18
19        if not postprocessing:
20            postprocessing = lambda x: x
21
22        exploss = None
23
24        for iter in range(start_iter + 1, iterations + 1):
25            # You might want to print the progress every few iterations.
26

```

```

27     # c, gin, gout = word2vecModel(
28     #     centerWord,
29     #     windowSize1,
30     #     context,
31     #     word2Ind,
32     #     centerWordVectors,
33     #     outsideVectors,
34     #     dataset,
35     #     word2vecLossAndGradient,
36     # )
37     # loss += c / batchsize
38     # grad[: int(N / 2), :] += gin / batchsize
39     # grad[int(N / 2) :, :] += gout / batchsize
40     loss, grad = f(x)
41     x -= step * grad
42
43     x = postprocessing(x)
44     if iter % PRINT_EVERY == 0:
45         if not exploss:
46             exploss = loss
47         else:
48             exploss = 0.95 * exploss + 0.05 * loss
49         print("iter %d: %f" % (iter, exploss))
50
51     if iter % SAVE_PARAMS_EVERY == 0 and useSaved:
52         save_params(iter, x)
53
54     if iter % ANNEAL_EVERY == 0:
55         step *= 0.5
56     return x

```

---

(c) (4 points) Show time! Now we are going to load some real data and train word vectors with everything you just implemented! We are going to use the Stanford Sentiment Treebank (SST) dataset to train word vectors, and later apply them to a simple sentiment analysis task. You will need to fetch the datasets first. To do this, run `sh get_datasets.sh`. There is no additional code to write for this part; just run `python run.py`.




Figure 2: The word vector distribution figure.